

✓ Fundamentals of Artificial Intelligence Programme (2024/25 Q1)

Multiobjective optimization for decision support

Dr. Jazmin Zatarain Salazar

In this assignment, we use two python libraries, one for exploratory modeling analysis [The EMA workbench](#) and one for multiobjective optimization [Project Platypus](#), you can install both with pip as follows:

```
!pip install ema_workbench platypus-opt
```

```

Requirement already satisfied: ema_workbench in /usr/local/lib/python3.10/dist-packages (2.5.2)
Requirement already satisfied: platypus-opt in /usr/local/lib/python3.10/dist-packages (1.3.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from ema_workbench) (1.26.4)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from ema_workbench) (2.1.4)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (from ema_workbench) (1.3.2)
Requirement already satisfied: salib>=1.4.6 in /usr/local/lib/python3.10/dist-packages (from ema_workbench) (1.5.1)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from ema_workbench) (3.7.1)
Requirement already satisfied: statsmodels in /usr/local/lib/python3.10/dist-packages (from ema_workbench) (0.14.2)
Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-packages (from ema_workbench) (0.13.1)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from ema_workbench) (4.66.5)
Requirement already satisfied: multiprocessing in /usr/local/lib/python3.10/dist-packages (from salib>=1.4.6->ema_workbench) (3.9.1)
Requirement already satisfied: scipy>=1.9.3 in /usr/local/lib/python3.10/dist-packages (from salib>=1.4.6->ema_workbench) (1.11.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->ema_workbench) (1.2.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->ema_workbench) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->ema_workbench) (4.52.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->ema_workbench) (1.4.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->ema_workbench) (24.0)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->ema_workbench) (10.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->ema_workbench) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib->ema_workbench) (2.9.0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->ema_workbench) (2023.3)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas->ema_workbench) (2024.1)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->ema_workbench) (1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->ema_workbench) (3.3.0)
Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.10/dist-packages (from statsmodels->ema_workbench) (0.5.6)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from patsy->statsmodels->ema_workbench) (1.16.0)
Requirement already satisfied: dill>=0.3.8 in /usr/local/lib/python3.10/dist-packages (from multiprocessing->salib>=1.4.6->ema_workbench) (0.3.8)

```

The Lake model

The goal of this assignment is to demonstrate the use of multiobjective evolutionary optimization, to learn how to visualize and interpret the optimization results, and to use performance metrics to assess the results obtained via multiobjective evolutionary optimization. We will use the lake problem as a test case, this is a classic problem initially developed by Carpenter et al. (1999) where the population of a city has to decide the amount of annual pollution it will release into a lake. In this exercise, we will use the adapted version in [Quinn et al. 2017](#) where the problem is defined as a state-based control problem (where actions are a function of the state of the system). In this case, the 'action' is the Phosphorous (P) emissions which are optimized to balance the economic benefits and the quality of the lake. Since this is a multi-objective problem, we need a flexible function to map the states to actions, so we use radial basis functions to parameterize the emission control policies. In fact, the MOEA will search for the optimal radii, centers and weights that yield good performance for the objectives of the lake model described below. See the paper for more details about the problem formulation.

The model is defined by the following equation:

$$X_{(t+1)} = X_t + a_t + \frac{(X_t^q)}{(1+X_t^q)} - bX_t + \epsilon_t$$

where X_t is the pollution at time t , a_t is the rate of anthropogenic pollution at time t , b is the lake's natural removal rate, q is the lake's natural recycling rate, ϵ_t is the rate of natural pollution at time t . The rate of anthropogenic pollution a_t is the decision variable where $a_t \in [0, 0.1]$.

There are four outcomes of interest. The first is the average concentration of phosphor in the lake.

$$f_{\text{phosphorus}} = \frac{1}{|T|} \sum_{t \in T} X_t$$

where $|T|$ is the cardinality of the set of points in time. The second objective is the economic benefit derived from polluting the lake defined as the discounted benefit of pollution minus the costs of having a polluted lake.

$$f_{\text{economic}} = \sum_{t \in T} \alpha a_t \delta^t$$

where α is the utility derived from polluting and δ is the discount rate. By default, α is 0.04. The third objective is related to the year over year change in the anthropogenic pollution rate.

$$f_{\text{inertia}} = \frac{1}{|T|-1} \sum_{t=1}^{|T|} I(|a_t - a_{t-1}| > \tau)$$

where I is an indicator function that is 0 if the statement is false, and 1 if the statement is true, τ is the threshold that is deemed undesirable, and is for illustrative purposes set to 0.2. Effectively, f_{inertia} is the fraction of years where the absolute value of the change in anthropogenic pollution is larger than τ . The fourth objective is the fraction of years where the pollution in the lake is below the critical threshold.

$$f_{\text{reliability}} = \frac{1}{|T|} \sum_{t \in T} I(X_t < X_{\text{crit}})$$

where I is an indicator function that is 0 if the statement is false, and 1 if the statement is true, X_{crit} is the critical threshold of pollution and is a function of both b and q .

The lake problem is characterized by both stochastic uncertainty and deep uncertainty. The stochastic uncertainty arises from the natural inflow. To reduce this stochastic uncertainty, multiple replications are performed and the average over the replication is taken. Deep uncertainty is presented by uncertainty about the mean μ and standard deviation σ of the lognormal distribution characterizing the natural inflow, the natural removal rate of the lake β , the natural recycling rate of the lake q , and the discount rate δ . The table below specifies the ranges for the deeply uncertain factors, as well as their best estimate or default values.

✓ Lake model implementation in python

```
1 import math
2 import numpy as np
3 from scipy.optimize import brentq
4
5
6 def get_antropogenic_release(xt, c1, c2, r1, r2, wl):
7     '''
8
9     Parameters
10    -----
11    xt : float
12           polution in lake at time t
13    c1 : float
14           center rbf 1
15    c2 : float
16           center rbf 2
17    r1 : float
18           radius rbf 1
19    r2 : float
20           radius rbf 2
21    wl : float
22           weight of rbf 1
23
```

```

24     Returns
25     -----
26     float
27
28     note:: w2 = 1 - w1
29
30     '''
31
32     rule = w1 * (abs(xt - c1) / r1) ** 3 + (1 - w1) * (abs(xt - c2) / r2) ** 3
33     at1 = max(rule, 0.01)
34     at = min(at1, 0.1)
35
36     return at
37
38
39 # def lake_model(b=0.42, q=2.0, mean=0.02,
40 #               stdev=0.001, delta=0.98, alpha=0.4,
41 #               nsamples=100, myears=100, c1=0.25,
42 #               c2=0.25, r1=0.5, r2=0.5,
43 #               w1=0.5, seed=123):
44
45 # adapted from the following request
46 def lake_model(b=0.42, q=2.0, mean=0.02,
47               stdev=0.001, delta=0.98, alpha=0.41,
48               nsamples=150, myears=100,
49               c1=0.25, c2=0.25, r1=0.5, r2=0.5,
50               w1=0.5, seed=123):
51     '''runs the lake model for nsamples stochastic realisation using
52     specified random seed.
53
54     Parameters
55     -----
56     b : float
57         decay rate for P in lake (0.42 = irreversible)
58     q : float
59         recycling exponent
60     mean : float
61         mean of natural inflows
62     stdev : float
63         standard deviation of natural inflows
64     delta : float
65         future utility discount rate
66     alpha : float
67         utility from pollution
68     nsamples : int, optional
69     myears : int, optional
70     c1 : float
71     c2 : float
72     r1 : float
73     r2 : float
74     w1 : float
75     seed : int, optional
76         seed for the random number generator
77
78     Returns
79     -----
80     tuple
81
82     '''
83     np.random.seed(seed)
84     Pcrit = brentq(lambda x: x ** q / (1 + x ** q) - b * x, 0.01, 1.5)
85
86     X = np.zeros((myears,))
87     average_daily_P = np.zeros((myears,))
88     reliability = 0.0
89     inertia = 0
90     utility = 0
91
92     for _ in range(nsamples):

```

```

93     X[0] = 0.0
94     decision = 0.1
95
96     decisions = np.zeros(myears, )
97     decisions[0] = decision
98
99     natural_inflows = np.random.lognormal(
100         math.log(mean ** 2 / math.sqrt(stdev ** 2 + mean ** 2)),
101         math.sqrt(math.log(1.0 + stdev ** 2 / mean ** 2)),
102         size=myyears)
103
104     for t in range(1, myyears):
105         # here we use the decision rule
106         decision = max(w1 * (abs(X[t - 1] - c1) / r1) ** 3 + (1 - w1) * (abs(X[t - 1]
107         decision = min(decision, 0.1)
108         decisions[t] = decision
109
110         X[t] = (1 - b) * X[t - 1] + X[t - 1] ** q / (1 + X[t - 1] ** q) + decision
111         average_daily_P[t] += X[t] / nsamples
112         if isinstance(average_daily_P[t], (bool, np.bool_)):
113             average_daily_P[t] = float(average_daily_P[t])
114         # print(f"average_daily_P[{t}]: {average_daily_P[t]}, type: {type(average_daily_P[t])}")
115
116
117         reliability += np.sum(X < Pcrit) / (nsamples * myyears)
118         inertia += np.sum(np.absolute(np.diff(decisions)
119                               < 0.02)) / (nsamples * myyears)
120         utility += np.sum(alpha * decisions * np.power(delta,
121                                                         np.arange(my
122     max_P = np.max(average_daily_P)
123     # if isinstance(max_P, (bool, np.bool_)):
124     #     max_P = 0.0 #
125
126     # print(f"max_P: {max_P}, type: {type(max_P)}")
127
128
129     return max_P, utility, inertia, reliability

```

✓ 1. Connecting the lake model with the EMA workbench.

Given the Python implementation of the lake problem above, adapt the code and connect it to the EMA workbench using the following lever ranges and uncertainty ranges:

Levers	Range	Default value
$r1$	0.0 – 2.0	0.5
$r2$	0.0 – 2.0	0.5.
$c1$	-2 – 2	0.25
$c2$	-2 – 2	0.25
$w1$	0.0-1.0	0.5
Uncertainties	Range	Default value
μ	0.01 – 0.05	0.02
σ	0.001 – 0.005	0.0017
b	0.1 – 0.45	0.42
q	2 – 4.5	2
δ	0.93 – 0.99	0.98

You can follow [this tutorial](#) for guidance.

The outcomes in the EMA workbench refers to the objectives of the problem, in this case we have four. 1) maximum Phosphorous (to be minimized) 2) utility (to be maximized) 3) inertia (to be maximized), and reliability (to be maximized). Use an alpha value of 0.41, with number of samples= 150, and number of years = 100.

```

1 # target:
2 # Maximum phosphorus concentration: should be minimized, because higher phosphorus concentration means deteri
3 # Economic benefits: should be maximized to achieve long-term economic benefits through pollution control s
4 # Inertia: should be maximized to maintain consistency and stability of pollution policies between years.
5 # Reliability: should be maximized to ensure that lake pollution is below the critical value in more yea
6
7
8 from ema_workbench import (Model, RealParameter, ScalarOutcome, MultiprocessingEvaluator)
9
10 lakemodel = Model('lakemodel', function=lake_model)
11
12 # Levers
13 lakemodel.levers = [RealParameter('r1', 0.0, 2.0),
14                     RealParameter('r2', 0.0, 2.0),
15                     RealParameter('c1', -2.0, 2.0),
16                     RealParameter('c2', -2.0, 2.0),
17                     RealParameter('w1', 0.0, 1.0)]
18
19 # Uncertainties
20 lakemodel.uncertainties = [RealParameter('mean', 0.01, 0.05),
21                             RealParameter('stdev', 0.001, 0.005),
22                             RealParameter('b', 0.1, 0.45),
23                             RealParameter('q', 2, 4.5),
24                             RealParameter('delta', 0.93, 0.99)]
25
26 # Outcomes
27 lakemodel.outcomes = [ScalarOutcome('max_P', kind=ScalarOutcome.MINIMIZE),
28                        ScalarOutcome('utility', kind=ScalarOutcome.MAXIMIZE),
29                        ScalarOutcome('inertia', kind=ScalarOutcome.MAXIMIZE),
30                        ScalarOutcome('reliability', kind=ScalarOutcome.MAXIMIZE)]

```

1 开始借助 AI 编写或生成代码。

2. How would you introduce a constrain within the optimization to reflect a desired performance threshold for a given objective?

You don't actually have to perform this step, simply specify how you would go about establishing a constraint in the optimization formulation, in such a way that it only finds solutions with maximum pollution (max phosphorous) of 0.85.

```

1 # add a constraint to the optimization problem
2 from ema_workbench import Constraint
3 constraints = [Constraint("max pollution", outcome_names="max_P", function=lambda x: max(0.85, x))]
4

```

3. Run the optimization and track the performance metrics


Tip: the EMA Workbench uses [Platypus](#) to run the optimization via the evaluator class, you can also collect metrics during runtime specifying the convergence option. Below is a sample snippet on how to run the optimization and collect performance metrics during runtime.

```

1 from ema_workbench import MultiprocessingEvaluator, ema_logging
2 from ema_workbench.em_framework.optimization import (HyperVolume,
3                                                       EpsilonProgress)
4
5
6 convergence_metrics = [HyperVolume(minimum=[0, 0, 0, 0], maximum=[1, 1.01, 1.01, 1.01]),
7                         EpsilonProgress()]

```

```
8
9
10 ema_logging.log_to_stderr(ema_logging.INFO)
11
12
13 with MultiprocessingEvaluator(lakemodel) as evaluator:
14     results, convergence = evaluator.optimize(
15         nfe=10000,
16         searchover='levers',
17         epsilons=[0.1, 0.1, 0.01, 0.1],
18         convergence=convergence_metrics,
19         constraints=constraints)
20
21
```

 [MainProcess/INFO] pool started with 2 workers

0%	0/10000 [00:00<?, ?it/s]
1% █	100/10000 [00:09<15:17, 10.79it/s]
2% █	200/10000 [00:17<14:17, 11.43it/s]
3% ██	300/10000 [00:25<13:33, 11.93it/s]
4% ██	400/10000 [00:34<13:53, 11.52it/s]
5% ███	500/10000 [00:43<13:48, 11.46it/s]
6% ███	600/10000 [00:51<13:25, 11.66it/s]
7% ███	700/10000 [01:00<13:32, 11.44it/s]
8% ███	800/10000 [01:10<13:50, 11.07it/s]
9% ████	900/10000 [01:17<12:41, 11.95it/s]
9% ████	900/10000 [01:28<12:41, 11.95it/s]
10% ████	1000/10000 [01:28<13:56, 10.76it/s]
11% ████	1100/10000 [01:36<13:04, 11.35it/s]
12% ████	1200/10000 [01:45<12:49, 11.44it/s]
13% ████	1300/10000 [01:52<12:16, 11.82it/s]
14% ████	1400/10000 [02:01<11:57, 11.99it/s]
15% ████	1500/10000 [02:09<11:50, 11.96it/s]
16% ████	1600/10000 [02:16<11:14, 12.45it/s]
17% ████	1700/10000 [02:26<11:46, 11.75it/s]
18% ████	1800/10000 [02:33<10:56, 12.49it/s]
19% ████	1900/10000 [02:42<11:15, 11.99it/s]
20% ████	2000/10000 [02:49<10:31, 12.67it/s]
21% ████	2100/10000 [02:58<10:51, 12.13it/s]
22% ████	2200/10000 [03:04<10:04, 12.91it/s]
23% ████	2300/10000 [03:14<10:35, 12.12it/s]
24% ████	2400/10000 [03:22<10:37, 11.93it/s]
25% ████	2500/10000 [03:31<10:23, 12.03it/s]

双击（或按回车键）即可修改

4. Selecting the objectives.

The outputs from the optimization runs will contain the decision variables (i.e. the parameters of the radial basis functions) and objectives combined. Create a data structure that only contains the objective values without the decision variables.

Tip: each row in the output matrix represents a different solution with it's obective values, the first columns are the decision variables and the last columns are the objective values.

1 results

	r1	r2	c1	c2	w1	max_P	utility	inertia	reliability
0	0.475550	0.976891	0.206675	0.183765	0.942480	0.097738	0.267078	0.990000	1.000000
1	0.085196	1.238544	0.240845	-0.261639	0.773910	0.522222	0.729131	0.987867	0.866733
2	0.085196	1.499870	0.240801	-0.405092	0.825415	0.838874	0.910375	0.985933	0.725267
3	0.085196	0.804180	0.242007	-0.012470	0.847954	0.481739	0.707849	0.743333	0.883867
4	0.086045	1.238544	0.240845	-0.261639	0.773910	0.356860	0.637278	0.989067	0.939400
5	0.456780	1.499182	0.327283	-0.298367	0.944601	0.194646	0.534778	0.990000	1.000000
6	0.085196	0.709425	0.240825	-0.062171	0.820034	0.650320	0.807812	0.986800	0.805200

1 convergence

	hypervolume	epsilon_progress	nfe
0	0.0	0	0
1	0.0	5	100
2	0.0	6	1100
3	0.0	6	2100
4	0.0	6	3100
5	0.0	6	4100
6	0.0	6	5100
7	0.0	10	6100
8	0.0	16	7100
9	0.0	24	8100
10	0.0	29	9100
11	0.0	31	10021

```

1 import pandas as pd
2
3 objective_columns = ['max_P', 'utility', 'inertia', 'reliability']
4
5
6 objectives_df = results[objective_columns]
7
8

```

```
9 print(objectives_df)
```

```
10
```



	max_P	utility	inertia	reliability
0	0.097738	0.267078	0.990000	1.000000
1	0.522222	0.729131	0.987867	0.866733
2	0.838874	0.910375	0.985933	0.725267
3	0.481739	0.707849	0.743333	0.883867
4	0.356860	0.637278	0.989067	0.939400
5	0.194646	0.534778	0.990000	1.000000
6	0.659329	0.807812	0.986800	0.805200

```
1 objectives_df.to_csv('objectives_only.csv', index=False)
```

```
2 results.to_csv('results.csv', index=False)
```


```
3 convergence.to_csv('convergence.csv', index=False)
```

✓ 5. Visualizing the results.

Present visually the results of the Pareto optimal solutions (in the objective space), feel free to be creative! Provide a brief discussion the results, are there any tradeoffs observed?

Tip: If you need inspiration check out the EMA workbench [parcoords](#).

```
1 from ema_workbench.analysis import parcoords
2 import matplotlib.pyplot as plt
3
4 data = objectives_df
5
6
7 limits = parcoords.get_limits(data)
8 limits.loc[0, ['max_P', 'utility', 'reliability']] = 0
9 limits.loc[0, ['inertia']] = 0.7
10 limits.loc[1, ['max_P', 'utility', 'inertia', 'reliability']] = 1
11
12
13 paraxes = parcoords.ParallelAxes(limits)
14 paraxes.plot(data)
15 paraxes.invert_axis('max_P')
16
17
18 plt.show()
```





Yes, the parallel coordinates plot reveals several trade-offs among the objectives in the multi-objective optimization problem. Specifically, it shows that reducing maximum phosphorus levels (Max_P) often compromises utility and reliability. Additionally, there's an observable trade-off between inertia and reliability, suggesting that efforts to maintain the status quo (inertia) may adversely affect the reliability of the system.

6. Establishing a performance threshold.

Show visually only the solutions from the Pareto set that yield a reliability above 80%, and briefly discuss the results.

```
1 import plotly.express as px
2
3
4 filtered_data = objectives_df[objectives_df['reliability'] > 0.8]
5
6
7 limits = parcoords.get_limits(data)
8 limits.loc[0, ['max_P', 'utility', 'reliability']] = 0
9 limits.loc[0, ['inertia']] = 0.7
10 limits.loc[1, ['max_P', 'utility', 'inertia', 'reliability']] = 1
11
12
13 paraxes = parcoords.ParallelAxes(limits)
14 paraxes.plot(filtered_data)
15 paraxes.invert_axis('max_P')
16
17
18 plt.show()
19
```



The plot illustrates the balance between Max Phosphorus, Utility, Inertia, and Reliability across different solutions. A trade-off exists between Max Phosphorus and Reliability, where lower phosphorus doesn't always ensure higher reliability. Similarly, higher utility often corresponds with lower inertia, meaning more flexible, dynamic solutions tend to perform better.

Solutions that achieve high reliability vary in their performance regarding phosphorus and utility, showing multiple paths to reliability. A cluster of solutions also balances high utility and reliability, suggesting an optimal compromise. This visualization helps decision-makers weigh trade-offs based on whether environmental impact or operational performance is prioritized.

✓ 7. Performance metrics

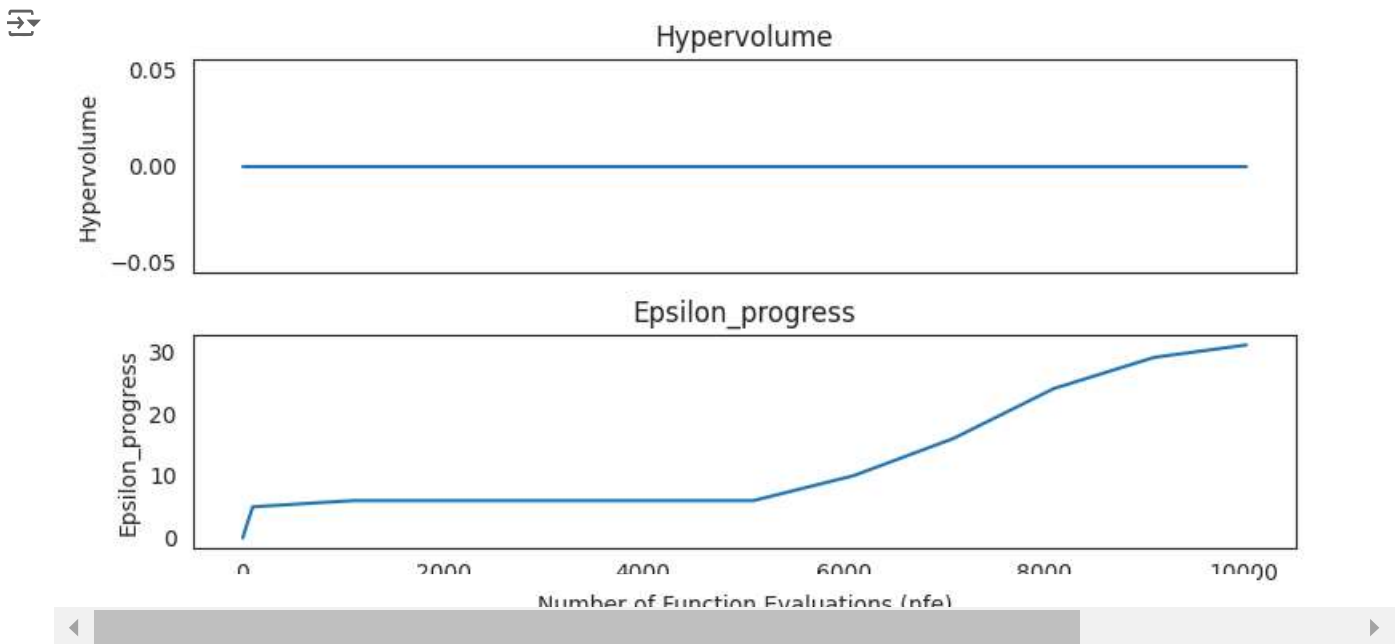
Show in a dataframe the results from the metrics (convergence) collected during the optimization. Plot the metrics (e.g hypervolume or epsilon progress) as a function of the number of function evaluations (nfe). Provide a brief discussion of the results.

```
1 import matplotlib.pyplot as plt
2 import seaborn as sns
3
4 # Set the style
5 sns.set_style("white")
6
7 # List of possible metrics we want to check if they exist in the DataFrame
8 possible_metrics = ['hypervolume', 'epsilon_progress', 'generational_distance', 'epsilon_indicator', 'inverted_gd',
9
10 # Filter out the metrics that exist in the DataFrame
11 available_metrics = [metric for metric in possible_metrics if metric in convergence.columns]
12
13 # Determine the number of plots based on available metrics
14 num_plots = len(available_metrics)
15
16 # Create subplots dynamically based on the number of available metrics
17 fig, axes = plt.subplots(nrows=num_plots, figsize=(8, 2 * num_plots), sharex=True)
18
19 # If only one metric available, wrap the axes object in a list
```

```

20 if num_plots == 1:
21     axes = [axes]
22
23 # Plotting each available metric
24 for ax, metric in zip(axes, available_metrics):
25     ax.plot(convergence['nfe'], convergence[metric], label=metric)
26     ax.set_title(metric.capitalize())
27     ax.set_ylabel(metric.capitalize())
28
29 # Set common x-axis label
30 axes[-1].set_xlabel('Number of Function Evaluations (nfe)')
31
32 # Adjust layout
33 plt.tight_layout()
34
35 # Show the plot
36 plt.show()
37

```



- ✓ Extra credit: Visualize the phosphorous release from the policy with the highest reliability. If you plot phosphorous release as a function of time, what do you observe?

```

1 def lake_model(b=0.42, q=2.0, mean=0.02, stdev=0.001, delta=0.98, alpha=0.41,
2               nsamples=150, myears=100, c1=0.25, c2=0.25, r1=0.5, r2=0.5, w1=0.5, seed=123):
3     """
4     Runs the lake model simulation for a given set of parameters and policies.
5
6     Parameters:
7     - b, q, mean, stdev, delta, alpha : float, environmental and economic parameters.
8     - nsamples : int, number of samples for stochastic realization.
9     - myears : int, number of years to simulate.
10    - c1, c2, r1, r2, w1 : float, parameters for the decision-making policy.
11    - seed : int, random seed for reproducibility.
12
13    Returns:
14    - tuple (float, float, float, float, np.array), containing max P, utility, inertia, reliability,
15      and phosphorous release over time for visualization.
16    """
17    np.random.seed(seed)
18    Pcrit = brentq(lambda x: x ** q / (1 + x ** q) - b * x, 0.01, 1.5)
19
20    X = np.zeros((myears,)) # Pollution levels
21    phosphorous_release_over_time = np.zeros((myears,)) # Store phosphorous release values for each year
22    reliability = 0.0

```

```

23     inertia = 0
24     utility = 0
25
26     for _ in range(nsamples):
27         X[0] = 0.0 # Start with no pollution
28         decision = 0.1 # Initial decision
29
30         decisions = np.zeros(myears,)
31         decisions[0] = decision
32         natural_inflows = np.random.lognormal(math.log(mean ** 2 / math.sqrt(stdev ** 2 + mean ** 2))
33                                                math.sqrt(math.log(1.0 + st
34                                                size=myyears))
35
36         for t in range(1, myears):
37             decision = get_antropogenic_release(X[t - 1], c1, c2, r1, r2, w1)
38             decisions[t] = decision
39             X[t] = (1 - b) * X[t - 1] + X[t - 1] ** q / (1 + X[t - 1] ** q) + decisio
40             phosphorous_release_over_time[t] += decision / nsamples # Accumulate average release
41
42             reliability += np.sum(X < Pcrit) / (nsamples * myears)
43             inertia += np.sum(np.absolute(np.diff(decisions) < 0.02)) / (nsamples * myears)
44             utility += np.sum(alpha * decisions * np.power(delta, np.arange(myears))) / nsamples
45
46     max_P = np.max(X)
47
48     return max_P, utility, inertia, reliability, phosphorous_release_over_time
49
50
51 1 # Identify the policy with the highest reliability, and it's not only one
52 max_reliability = results['reliability'].max()
53 highest_reliability_policies = results[results['reliability'] == max_reliability]
54
55 2 plt.figure(figsize=(12, 6))
56
57 3 for _, policy in highest_reliability_policies.iterrows():
58     max_P, utility, inertia, reliability, phosphorous_release = lake_model(
59         c1=policy['c1'],
60         c2=policy['c2'],
61         r1=policy['r1'],
62         r2=policy['r2'],
63         w1=policy['w1']
64     )
65
66     plt.plot(np.arange(100), phosphorous_release, label=f'Policy: r1={policy["r1"]}, r2={policy["r2"]}, w1={poli
67
68 4 plt.xlabel('Year')
69 plt.ylabel('Phosphorus Release')
70 plt.title('Phosphorus Release Over Time for Policies with Highest Reliability')
71 plt.legend()
72 plt.grid(True)
73 plt.show()

```

